Revisiting Difficulty Control for Blockchain Systems

No Author Given

No Institute Given

Abstract. The Bitcoin whitepaper [1] states that the security of the system is guaranteed as long as honest miners control more than half of the current total computational power. The whitepaper assumes a static difficulty, thus it is equally hard to solve a cryptographic proof-of-work puzzle for any given moment of the system history. However, the real Bitcoin network is using an adaptive difficulty adjustment mechanism.

In this paper we introduce and analyze a new kind of attack on a mining difficulty retargeting function used in Bitcoin. A malicious miner is increasing his mining profits from the attack, named coin-hopping attack, and, as a side effent, an average delay between blocks is increasing.

We propose an alternative difficulty adjustment algorithm in order to reduce an incentive to perform coin-hopping, and also to improve stability of inter-block delays. Finally, we evaluate the presented approach and show that the novel approach performs better than the original algorithm of Bitcoin.

1 Introduction

Blockchain systems have attracted significant amount of interest after the Bitcoin whitepaper [1] was published in 2008. Bitcoin security relies on a distributed protocol which maintains a distributed ledger. In the protocol miners are trying to find a partial hash collision in order to generate a valid block by iterating over nonce field values.

Alternative systems may rely on other types of computational puzzles rather than finding a partial hash collision, e.g., [2,3]. Nevertheless, all of them assume some algorithm that changes the difficulty of the puzzle dynamically. An algorithm for difficulty readjustment is required in order to make an open blockchain system working stable in the face of participants joining and leaving the system (resulting in constantly changing available computational power for solving the puzzles), and also to stabilize mean latency between blocks.

The difficulty readjustment algorithm in Bitcoin assumes that the total computational power involved in the mining process does not significantly change from epoch to epoch. In contrast, real networks show that a significant variance in computational power happens over long periods. For example, we show in this paper that due to continuous growth of computational power in the Bitcoin network a mean delay between blocks differs from an expected value by 7%.

Noteworthy, exponential growth of computational power, often observed in practice, is the absolutely worst case (regarding the mean block delay divergence) for the Bitcoin's difficulty readjustment algorithm [4].

In this paper we also consider a new type of miner behavior with regards to difficulty readjustment which provides unfair advantage to the miner, and also makes inter-block delays worse. We call the discovered strategy the coin-hopping attack following the "pool-hopping" term raised in [5]. In this attack, an adversarial miner is switching from mining one coin to another in the beginning of an epoch, then he is switching back in the beginning of next epoch when difficulty becomes lower. We show how adversarial mining profit is increasing for Bitcoin's difficulty readjustment function, and how inter-block delay suffers from the coin-hopping attack.

As a solution for the significant variance in computational power and also in order to reduce incentive of the described coin-hopping strategy, we propose an alternative difficulty readjustment procedure. We show that the proposed solution is better suited for exponential growth of the total mining power. It also reduces profit and negative side-effects of the coin-hopping attack.

1.1 Related Work

In this section we provide an overview of known formal and informal studies with regard to the dynamic nature of the difficulty parameters in Bitcoin. Following the well known paper of Garay et al. [6], generalizing the Bitcoin backbone protocol in a static difficulty setting, a newer paper from the same authors [7] is providing a positive answer on whether basic security properties of the Bitcoin backbone protocol (common prefix, chain quality and chain growth) hold in case of dynamic difficulty, in a cryptographic setting with an arbitrary adversary. Nevertheless, studying concrete attacks against the real protocol is still needed.

The Timejacking attack [8] allows an attacker to first shift the network time at a victims node (which is calculating network time by averaging timestamps it gets regularly from neighbors) and then force the victim node to reject a block with a specially crafted timestamp (other nodes are accepting). The time wrapping attack [9] is exploiting the fact that Bitcoin is using difference in timestamps between last and first block in an epoch, instead of the last block in an epoch and last block in a previous epoch. By using specially crafted timestamps for the last block of each epoch, an attacker can produce more blocks for a time window with more work contributed to his chain. The difficulty raising attack, introduced in [10], allows an attacker to discard n-depth block, for any n, and for any computational power of the attacker, with probability 1 if he is willing to wait long enough.

Another paper [4] is introducing an alternative difficulty readjustment function designed to work better than Bitcoin's not just for almost constant mining power but also when the power is growing exponentially with time. We provide comparison with this function in the paper in one of the following sections.

The paper is organized as follows: in Section 2 we provide a detailed view of Bitcoin's readjustment function. In Section 3 we introduce the coin-hopping

attack, followed by the definition of an improved difficulty readjustment function described in section 4. Section 5 provides results of evaluation of the presented approach.

2   Bitcoin Mining

The concept of Bitcoin mining was introduced in Section 4 of the Bitcoin whitepaper [1], and then discussed in detail in the papers [4],[7]. Bitcoin miner generates a block by iterating over a *nonce* value and calculating the hash of a block with the nonce value included. For a block $\mathcal{B}$ to be valid, a value of a hash function has to be less than the current *target* T, $hash(\mathcal{B}) < T$, where $hash$ is an ideal cryptographic hash function. Hardness to find a block could be expressed also via *difficulty* D as $D = \frac{1}{T}$. If output of the *hash* function is μ bits long then the probability to generate a block by doing q requests to the hash function is $\frac{T \cdot q}{2^\mu} = \frac{q}{D \cdot 2^\mu}$. We define miner's *hashrate* R as $R = \frac{q_s}{2^\mu}$, where q_s is number of queries done by miner s per time unit. The probability to generate a block within a time unit is then $\frac{R}{D}$.

Every M blocks ($M = 2016$ for Bitcoin) the difficulty is recalculated as

$$D_{i+1} = D_i \cdot \frac{M \cdot |\Delta|}{S_m} \tag{1}$$

where $|\Delta|$ is the expected time interval between blocks and S_m is the actual time spent to generate M blocks. For the Bitcoin network, the observed time interval of $\approx$*9 minutes 20 seconds* is less than the planned value of $|\Delta| = 10$ *minutes* due to continuous growth of the computational power of the network. Difficulty recalculation interval $M = 2016$ has been chosen to recalculate difficulty every 2 weeks on average. The epoch length is big enough to see the computational power of the network being changing over it: mean delay is close to the planned 10 minutes right after target recalculation, whereas at the end of an epoch it is less than 9 minutes in average.

The next section describes an attack that allows to gain benefits by attacking the recalculation algorithm.

3   Coin-hopping Attack

We consider the following attack involving an adversarial miner $\mathcal{A}$:

- There are at least 2 possible coins (C_1,C_2) $\mathcal{A}$ can contribute to. Without a loss of generality, we assume that each of them provides about the same profitability of the mining activity.
- $\mathcal{A}$ is mining coin C_2 before the beginning of an epoch A. At the beginning of A he switches to mine coin C_1.
- Without the contribution of miner $\mathcal{A}$ the total mining power of the C_2 network for the epoch decreases.

– For an epoch B right after epoch A, the difficulty of C_2 gets readjusted to a lower value. So $\mathcal{A}$ starts mining C_2 again with a lower difficulty.

We call this strategy a *coin-hopping attack*.

To calculate the profit the adversarial miner gains from this attack, we use Bitcoins' difficulty recalculation function and assume a constant network hashrate (with respect to the rest of the network, without our adversarial miner). We denote the hashrate of miners not participating in the coin-hopping attack as R_0 in both C_1 and C_2, and we denote the hashrate of the adversarial miner as $R_a = R_0 \cdot p, 0 < p < 1$. Before epoch A the adversary is mining coin C_2, thus the difficulty of the C_2 network is $D_0 = (R_0 + R_a) \cdot |\Delta|$ (see Section 3.1 in [4]). During the epoch A the difficulty of the C_2 network is still D_0, and $\mathcal{A}$ switches to mine coin C_1 at a difficulty $D_1 = R_0 \cdot |\Delta|$ calculated from honest miners hashrate R_0 only. During the epoch B the adversary starts mining of C_2, now at difficulty D_1, while honest miners on chain C_1 continue to mine it with higher difficulty D_0. After that $\mathcal{A}$ continues to switch between chains C_1 and C_2 always mining on the chain with lower difficulty D_1, spending $R_0 \cdot |\Delta|$ computational power per block, whereas honest miners spend $(R_0 + R_a)|\Delta|$ computational power per block.

Every epoch honest miners with hashrate R_0 will generate $\frac{M \cdot R_0}{R_0 + R_a}$, blocks, whereas $\mathcal{A}$ will generate $\frac{M \cdot R_a}{R_0}$ blocks. If $\mathcal{W}$ is block reward, the additional profit of the adversary is calculated as the difference of what he mines based on the lower difficulty in contrast to the difficulty he would mine at without hopping between the coins:

$$
\begin{aligned}
\mathcal{W} \cdot M \cdot \frac{R_a}{R_0} - \mathcal{W} \cdot M \cdot \frac{R_a}{R_0 + R_a} &= \\
= \mathcal{W} \cdot M \cdot (\frac{R_a}{R_0} - \frac{R_a}{R_0 + R_a}) &= \\
= \mathcal{W} \cdot M \cdot \frac{R_a \cdot (R_0 + R_a) + R_a \cdot R_0}{R_0 \cdot (R_0 + R_a)} &= \\
= \mathcal{W} \cdot M \cdot \frac{R_a^2}{R_0 \cdot (R_a + R_0)} &= \\
= \mathcal{W} \cdot M \cdot \frac{R_0^2 \cdot p^2}{R_0 \cdot (R_0 \cdot p + R_0)} &= \\
= \mathcal{W} \cdot M \cdot \frac{p^2}{1 + p}
\end{aligned}
\tag{2}
$$

Remarkably, under such an attack the mean time between blocks in both chains C_1 and C_2 will be

$$
T_a = \frac{T}{2}\left(\frac{R_0 + R_a}{R_0} + \frac{R_0}{R_0 + R_a}\right) = T\left(1 + \frac{p^2}{2(1 + p)}\right)
\tag{3}
$$

which is bigger than the planned time T.

The next section provides a new algorithm for difficulty adjustment.

4 Improved Difficulty Adjustment

The difficulty adjustment algorithm employed by Bitcoin works as designed: if the hash rate of the network is constant, it yields to the desired block rate. However it does not achieve the desired block rate in other situations and is vulnerable to the attack described in 3. In this section we are going to propose an alternative difficulty adjustment algorithm.

First, we state properties of an ideal difficulty update algorithm:

1. It should be resistant to known types of attacks based on difficulty manipulation.
2. It should lead to an almost constant desired block rate for random fluctuations in the network hashrate.

We propose a difficulty adjustment algorithm based on the well-known linear least squares method[11], we name it *linear algorithm*. In the simplest case of pair linear regression $y = kx + b$, coefficients may be calculated as follows:

$$\begin{cases} k = \frac{\overline{xy} - \overline{x}\overline{y}}{\overline{x^2} - \overline{x}^2} \\ b = \overline{y} - k\overline{x} \end{cases} \tag{4}$$

Note that for accurate difficulty prediction we should use a few last observed difficulties, rather than just one, as implemented in Bitcoin, but it is still possible to use this algorithm right after second epoch.

The next section provides an evauation of the presented approach.

5 Evaluation

We now present simulation results that show that the method proposed in Section 4 outperforms Bitcoin's difficulty update algorithm. We will regard difficulty growth in this section, keeping in mind the fact, that it is closely related with network hash rate, which is usually considered in literature.

5.1 Experiment

First, we observe exponential difficulty growth as it occures in practice in the Bitcoin network. As we already mentioned, exponential difficulty growth is the absolutely worst case possible for Bitcoin's difficulty retargeting algorithm. For simplicity we regard a situation, when hash rate increases by 10% each epoch, more complicated research of exponential difficulty growth can be found in [4]. Figure 1 presents difficulty function over epochs, which is 2016 blocks in Bitcoin.

Note that the difficulty calculated from Bitcoin algorithm is always significantly lower than the real one. This leads to a *9 min 5 sec* time interval between blocks, which is ≈10% lower then desired *10 min* interval. Difficulty, calculated by the linear algorithm is constantly lower, then the real one, but still much

Fig. 1. Real difficulty (red) and difficulties calculated from bitcoin (black) and linear (blue) algorithms in situation of exponential hash rate growth

closer to it. Mean time interval between blocks in this approach is *9 min 45 sec*, which is much closer to the desired one.

Furthermore, we looked at the error in percentage of the difficulty. That shows that the algorithm currently used in the Bitcoin network has a pretty constant average error of about 9.1%, while our algorithm only has a constant error of about 1.9%.

While the difficulty update algorithm, proposed in [4] leads to much better results for exponential difficulty growth with a constant rate, we should note, that our algorithm is much simpler and may be implemented with integer arithmetic only. Moreover, exponential difficulty growth is the simplification of the difficulty growth law, and it may be incorrect to expect it in some situations.

5.2   Real World Data

In order to further evaluate the presented approach, we also used real world data from the Bitcoin network. For this we compared the real world data of the difficulty calculated by the algorithm used in Bitcoin in comparison to the difficulty as it should have been. Furthermore, we do the same calculation based on our model. Results are similar as for the experimental data described in 5.1.

We further evaluated the average of the percentage of the differences, again both for the Bitcoin algorithm as well as for our approach and have found that in average, Bitcoins' algorithm has an error of about 12.3% while our approach only has an error of about 8.4%. This shows that our approach performs about 33% better than the approach currently used in the Bitcoin network.

5.3 Coin-hopping Attack

We consider a situation as described in the Section 3: an attacker with computational power R_a (for simplicity we suppose $R_a = 0.2 \cdot R_0$ in this section) turn on and turn off his mining to manipulate difficulty and minimize computational power, expended for block mining. Figure 2 represents difficulty as the function over epochs for this situation.
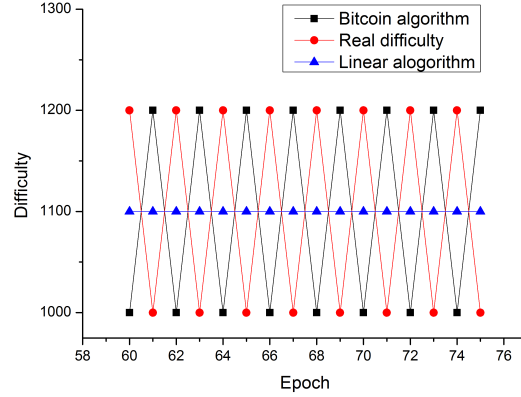
Fig. 2. Real difficulty (red) and difficulties calculated from bitcoin (black) and linear (blue) algorithms in situation of coin-hopping attack

Note that the difficulty calculated with the Bitcoin algorithm is always in antiphase with the real one and the attacker invests his computational power only when difficulty is low. The Bitcoin difficulty update algorithm leads to *10 min 10 sec* mean delay between blocks, which is in good correlation with 3. The linear algorithm also leads to enlarged time interval between blocks equal to *10 min 5 sec*, resulting in a two times lower deviation from the desired time. Obviously, the profit of the attacker is then also 2 times lower while using the the linear difficulty update algorithm, providing a servious improvement.

Thus, linear difficulty control algorithm, proposed in Section 4 is better than the one used in Bitcoin for the given situation, both in terms of block rate and in terms of attacker profit.

References

1. S. Nakamoto, Bitcoin: A Peer-to-Peer Electronic Cash System (2008). `arXiv:43543534534v343453`, `doi:10.1007/s10838-008-9062-0`.
   URL `http://s.kwma.kr/pdf/Bitcoin/bitcoin.pdf`

2. A. Miller, A. Juels, E. Shi, B. Parno, J. Katz, Permacoin: Repurposing bitcoin work for data preservation, in: Security and Privacy (SP), 2014 IEEE Symposium on, IEEE, 2014, pp. 475–490.

3. A. Biryukov, D. Khovratovich, Equihash: Asymmetric proof-of-work based on the generalized birthday problem, Ledger 2.

4. D. Kraft, Difficulty control for blockchain-based consensus systems, Peer-to-Peer Networking and Applications (2015) 1–17.

5. M. Rosenfeld, Analysis of bitcoin pooled mining reward systems, arXiv preprint arXiv:1112.4980.

6. J. Garay, A. Kiayias, N. Leonardos, The bitcoin backbone protocol: Analysis and applications, in: Advances in Cryptology-EUROCRYPT 2015, Springer, 2015, pp. 281–310.

7. J. A. Garay, A. Kiayias, N. Leonardos, The bitcoin backbone protocol with chains of variable difficulty.

8. The timejacking attack.
   URL http://culubas.blogspot.com

9. ArtForz, The time wrapping attack.
   URL https://bitcointalk.org/index.php?topic=43692.msg521772#msg521772

10. L. Bahack, Theoretical bitcoin attacks with less than half of the computational power, arXiv preprint arXiv:1312.7013.

11. C. L. Lawson, R. J. Hanson, Solving least squares problems, Vol. 161, SIAM, 1974.